

A Reproducible Pipeline for Streaming Log Analytics at Scale

Authors: Ravi Kumar; Anna Fischer

Affiliations: Indian Institute of Technology Delhi / Technical University of Munich

Correspondence: rkumar@example.org

Document: Typeset galley (published version)

Abstract

Operational logs are a primary signal for diagnosing distributed systems, yet pipelines that process them are often ad hoc and hard to reproduce.

We present a modular pipeline that combines schema-on-read ingestion with incremental aggregation, and evaluate it on three public workloads.

Results indicate consistent latency under sustained load and a clear path to horizontal scaling.

Keywords: stream processing, observability, reproducibility, distributed systems

1. Introduction

This manuscript is part of a realistic but synthetic Open Journal Systems instance used for non-production testing. The content is fictional; any resemblance to real studies is coincidental.

2. Methods

The study design, materials, and analysis described here are illustrative placeholders that mirror the structure of a genuine research article so that downstream tooling sees plausible inputs.

References

- [1] Kumar, R. (2023). Working notes. Indian Institute of Technology Delhi.
- [2] Fischer, A. (2023). Working notes. Technical University of Munich.