

Tail-Latency Aware Load Balancing for Microservice Meshes

Authors: Olof Kvist; Thao Tran

Affiliations: KTH Royal Institute of Technology / Vietnam National University

Correspondence: okvist@example.org

Document: Author manuscript (as submitted)

Abstract

Mesh sidecars make per-request routing decisions cheap, opening the door to latency-aware policies.

We propose a balancer that estimates downstream queue depth from recent response times and show reduced p99 latency without throughput loss.

Keywords: microservices, load balancing, tail latency, service mesh

1. Introduction

This manuscript is part of a realistic but synthetic Open Journal Systems instance used for non-production testing. The content is fictional; any resemblance to real studies is coincidental.

2. Methods

The study design, materials, and analysis described here are illustrative placeholders that mirror the structure of a genuine research article so that downstream tooling sees plausible inputs.

References

- [1] Kvist, O. (2024). Working notes. KTH Royal Institute of Technology.
- [2] Tran, T. (2024). Working notes. Vietnam National University.